LEGEND

DATA | BUSINESS/ PRODUCT | AI APPLICATION

GNFR INDIRECT SPEND ROADMAP

FY2026
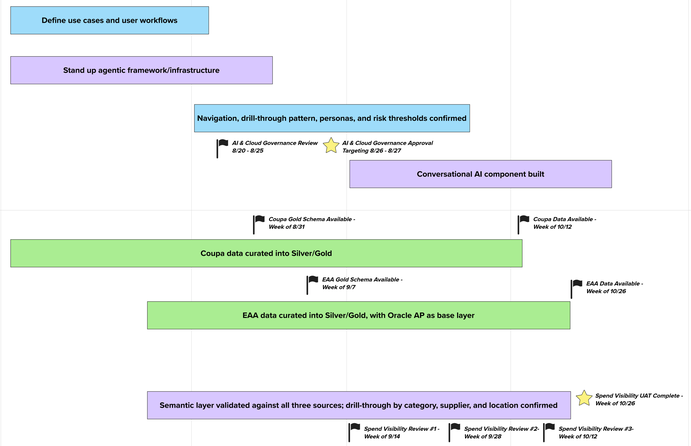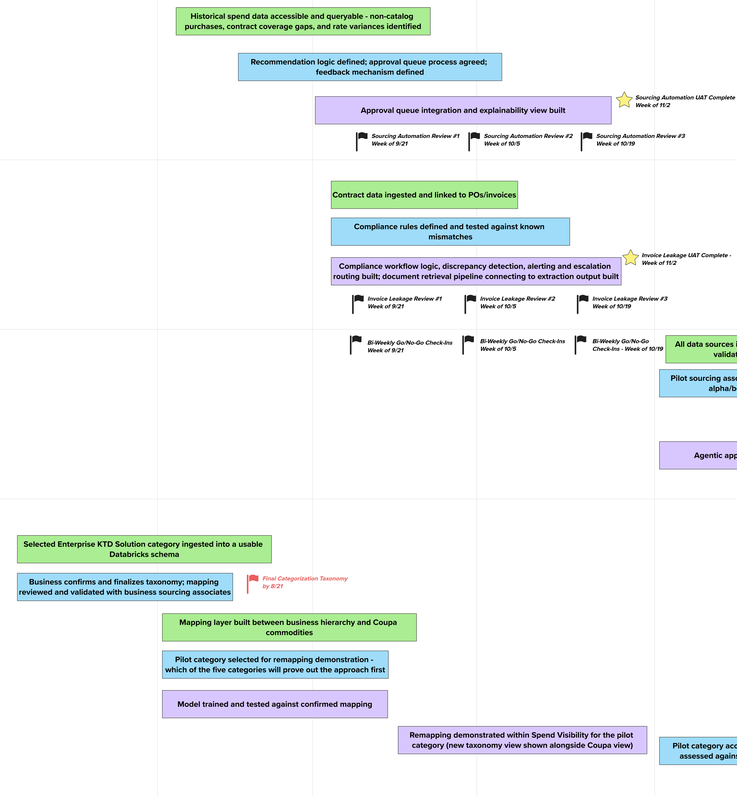
#	OBJECTIVE	DESCRIPTION	DEPENDENCIES/ RISKS	CURRENT (7/19 - 8/15)	P8 (8/16 - 9/12)	P9 (9/13 - 10/10)	P10 (10/11 - 11/7)	P11 (11/8 - 12/5)	P12 (12/6 - 1/2)	P13 (*Post-MVP*)
				Q2	**Q3**			**Q4**		
				P7 (7/19 - 8/15)						
1	OVERALL EXPERIENCE	• Web app delivered, integrating with Coupa (specific touchpoints still being defined) • Conversational AI component built for querying recommendations, evidence, and data directly • User personas for design confirmed • Navigation and drill-through pattern confirmed - single shared view for all users, manual filter/drill up-down, modeled on ontology-style navigation • Risk-tier thresholds (high/medium/low) defined using pre-defined, transparent business rules - shared standard applied across both Invoice Leakage and Sourcing Automation flagging • Data refresh frequency defined and surfaced to users • Configurable business rules (e.g. categorization confidence threshold, risk-tier cutoffs) adjustable without requiring redeployment	• AI Governance and Cloud Governance approvals must be completed; otherwise, blocks any AI capability from going live							
2	SPEND VISIBILITY	• Oracle AP curated as the base layer for 100% of GNFR spend, augmented by Coupa and EAA detail where available • Categorization taxonomy finalized (Coupa commodity hierarchy vs. alternative approach, e.g. UN/SPSC) and spend assigned to correct Sourcing Associate accordingly • Semantic layer returns accurate spend totals with drill-through by category, supplier, and location; validated against source data • MVP definition and end users/workflows confirmed ***- Complete*** • Agents can successfully query semantic layer and return spend insights	• KTD internal build requires Silver/Gold layers - additional effort needed vs. Palantir's Bronze-only approach • Coupa's category hierarchy remains an option for MVP, though known gaps noted, with UNSPSC codes suggested as an alternative							
3	DOCUMENT DIGITIZATION	• Invoice attachments accessible in Unity Catalog and feeding OCR pipeline • Contract attachment extraction feasibility evaluated (master agreements, amendments, SOWs) • OCR output structured, mapped into ontology, and queryable by AI agents • Line item detail unlocked and validated against sample of physical invoices/contract attachments • Invoice extraction approach decided (OCR vs. LLM-based extraction) • Quality bar defined for extraction accuracy before feeding downstream capabilities • Finance and Accounting buy-in secured; Marcy aligned on Coupa process changes	• POC needed to confirm value of invoice digitization and technical approach for Excel attachments • Contract extraction is new territory - no existing precedent anywhere, including Palantir • New contracting tool in flight - confirm scope/timeline before building contract parsing capability							
4	SOURCING AUTOMATION	• Sourcing recommendation logic tested and generating accurate, actionable savings opportunities • Contract leakage and supplier consolidation opportunities surfaced with quantified savings • Recommendation dismissal/feedback mechanism built, preventing resurfacing of false positives • Non-catalog spend patterns surfaced and attributable to correct Sourcing Associate • Conversational interface live and accessible to category managers • Approval queue built, tested, and process agreed upon with procurement team • Explainability view showing underlying data behind each recommendation, with a feedback loop to refine future logic	• Coupa test environment SAML/SSO not configured - blocks non-prod testing of discrepancy flagging							
5	INVOICE LEAKAGE	• Contract data ingested and linked to POs and invoices in Unity Catalog • Compliance rules (rate variance, contract max-value exposure, expiration/auto-renewal) producing accurate discrepancy flags; contract max-value applicability validated per category, since not universally used (e.g. services) • Discrepancy review queue live and routing correctly to Sourcing Associates • Compliance thresholds and escalation logic agreed with procurement team • Risk classification (high/medium/low) based on pre-defined, transparent business rules - validated with procurement team, not AI-determined • Contract term/SLA breach detection enabled, dependent on Document Digitization's contract extraction	• Contract maximum value ("not-to-exceed") field may not be broadly applicable across all categories, particularly services - needs validation with category teams once formed • SLA/contract-term breach detection depends on Document Digitization's contract extraction - scope risk if that capability slips • Coupa test environment SAML/SSO not configured - blocks non-prod testing of discrepancy flagging							
6	INDIRECT SPEND PARITY	• Spend Visibility, Document Digitization, Sourcing Automation, and Invoice Leakage tested end-to-end and validated against parallel build outputs • Parallel build offramp executed; knowledge transfer for in-scope capabilities completed • DevOps and support infrastructure in place for delivered capabilities • Conversational AI and recommendation explainability validated as at-parity with parallel build's chat and model-feedback experience • Recommendation and calculation outputs written back to Databricks Gold layer for enterprise reuse								
7	BUSINESS CATEGORY TAXONOMY MAPPING	• Business's category/desk hierarchy (currently in Epic, ~90% complete per business) confirmed and finalized across all five categories • Mapping layer built connecting business's hierarchy to Coupa commodity structure • Model trained and tested against confirmed mapping • **Remapping demonstrated within Spend Visibility for one pilot category (e.g. Professional Services)**, validated against the parallel build's approach for that same category, ahead of the November go/no-go • **Remaining four categories continue training in parallel**, rolling into production in phases post-MVP as each independently clears the agreed accuracy threshold **Note:** Coupa commodity hierarchy remains the taxonomy for Sourcing Automation and Invoice Leakage - unaffected by this workstream. Spend Visibility pilots the new taxonomy on one category to prove the remapping capability works, in parallel with the parallel build doing the same on that category, keeping the comparison valid. Full business-aligned categorization across all five categories remains targeted for early FY27.	• Need to confirm the parallel build is piloting the same single category on the same timeline, or the apples-to-apples comparison breaks • "Pilot-ready" accuracy/completeness threshold for the one demonstrated category not yet defined - needed before the Nov assessment