

User Story 2.2: Flagging Low-Confidence Categorization

As a Sourcing Associate, I want spend that can't be confidently categorized to be flagged rather than silently miscategorized, so I know which numbers to trust.

Description:

Surface a visible flag on spend records where categorization confidence is low, rather than silently assigning a best-guess category with no indication of uncertainty.

Narrative:

Not all spend will map cleanly to a category, regardless of which taxonomy is ultimately used (Coupa or the new business taxonomy). Rather than forcing every record into a category with equal apparent confidence, associates need to know when a categorization is uncertain so they don't unknowingly act on or report against a number that may be wrong. Confidence flagging should be treated as adjustable and improvable over time, not a hardcoded one-time judgment.

Business Value:

Protects trust in the categorization data as a whole - a single silently-miscategorized high-value transaction can undermine confidence in the entire dataset. Flagging low-confidence records lets associates apply judgment where the model can't yet be trusted, rather than assuming every category assignment is equally reliable.

Acceptance Criteria:

- Given a spend record's categorization confidence falls below a defined threshold, when displayed, then it is visibly flagged as low-confidence
 - **Threshold TBD**
- Given a confidence threshold, when it is adjusted (e.g. by an admin or business rule), then flagging behavior updates accordingly without requiring a redeploy
- Given a low-confidence record, when a user reviews it, then the reason for low confidence (e.g. no clean taxonomy match, conflicting signals) is available, not just a binary flag

Out of Scope:

- Automatically correcting or re-routing low-confidence records - this story covers flagging only, not remediation workflow