

M2: Document Digitization Requirements

Scope confirmed in session: Coupa attachments only (not EAA or other systems). Attachment types limited to invoices and contracts - not POs or other attachment types. Contract OCR is net-new territory (no prior attempt anywhere, including Palantir).

Epic 1: Coupa Invoice OCR Extraction

Description: Extract structured line-item detail from Coupa invoice attachments to close the categorization and compliance-checking gap left by summarized/truncated OCR today.

Backed by: `invoices_v1` (Bronze), `Invoice`, `LineItemEvidence`, `record.invoice-lines[]`

Known gap: Existing OCR tool summarizes rather than preserves structure (e.g. a 100-line invoice becomes a single line with a total). Attachment content can also differ from structured invoice line items in Coupa - not a guaranteed apples-to-apples source.

User Stories

- As a Data/AI Engineer, I want to extract line-item level detail (ex. item, quantity, unit price) from Coupa invoice attachments, so downstream categorization and compliance checks have granularity that the summarized OCR output doesn't provide today.
- As a Sourcing Associate, I want extracted invoice line items linked back to the parent invoice record, so I can verify extraction accuracy against the source document.
- As a Data/AI Engineer, I want to flag invoices where OCR-extracted content conflicts with or can't be reconciled against structured Coupa invoice data, so downstream consumers know when to trust the OCR result versus the source system record.

Epic 2: Coupa Contract OCR Extraction

Description: Extract structured terms from Coupa contract attachments (master agreements, amendments/appeals, and child documents like SOWs) to support invoice leakage detection and contract lifecycle monitoring. This is new capability with no existing precedent.

Backed by: `contracts_v1` (Bronze), `Contract`, `record.contract-clauses`, `record.contract-terms`, `record.contract-parties`

User Stories

- As a Data/AI Engineer, I want to extract pricing terms (contracted rate) from contract attachments, so invoice amounts can be checked against what was actually agreed rather

than relying only on structured contract table fields.

- As a Data/AI Engineer, I want to extract contract maximum value from contract attachments where not captured in structured fields, so leakage detection can flag spend approaching or exceeding that ceiling.
- As a Data/AI Engineer, I want to extract SLA terms (e.g. delivery timelines, performance commitments) from contract attachments, so term-breach detection has language to compare actual supplier performance against.
 - Payment terms, duration terms, etc.
- As a Sourcing Associate, I want extracted contract terms linked to the correct parent contract and any child documents (SOWs, amendments), so I see the full picture of what governs a given supplier relationship.
- ***Expiration or effective date***
- As a Product/Design lead, I want contract OCR treated as a discovery/POC effort before committing to a specific extraction approach, so we don't over-invest in a method before validating feasibility on real Kroger contracts.

Epic 3: OCR Output Quality and Structure Preservation

Description: Ensure OCR output preserves document structure well enough to support line-item-level use cases, rather than the summarized output the current tool produces.

Backed by: `DataQualityIssue` , `gold_source_freshness_metric`

User Stories

- As a Data/AI Engineer, I want a defined quality bar for OCR extraction (e.g. percentage of line items correctly parsed) before invoice/contract OCR output feeds downstream use cases, so Sourcing Automation and Invoice Leakage aren't built on unreliable extraction.
- As a Sourcing Associate, I want to know when an invoice or contract's OCR extraction is low-confidence or incomplete, so I don't act on missing or wrong detail.

Cross-Epic Open Items

- Confirm insurance-expiration tracking system of record (Carlos to check) - may become a future OCR/contract-monitoring use case, not yet MVP scope.
- Contract OCR quality bar and extraction approach not yet defined - genuinely new ground; needs a proof-of-concept pass before full scoping.