

👁️ M1: Spend Visibility Requirements

Note: This is a first pass to structure discovery into epics/stories. Coupa is the only source with a validated Bronze-to-Gold object map today. Per the Aug 3 session, Oracle AP is the primary/base data source for Spend Visibility (100% of GNFR spend), with Coupa and EAA augmenting it, but no equivalent object/relationship map exists yet for Oracle AP or EAA. Flagging that as an open dependency under each epic where it applies.

Epic 1: Unified Spend Foundation Across Oracle AP, Coupa, and EAA

Description: Establish Oracle AP as the base layer for 100% GNFR spend visibility, augmented by Coupa and EAA detail where available, so users always see a complete spend picture even where source-system detail is missing.

Backed by (Coupa side): SpendFact , gold_supplier_metric , po_spend , invoice_spend

Open dependency: Equivalent Bronze/Silver/Gold object map needed for Oracle AP (base layer) and EAA - not yet documented.

User Stories

- As a Sourcing Associate, I want to see total GNFR spend by supplier, category, and location (division, manufacturing plant, distribution center level) even when a transaction only exists in Oracle AP, so I have a baseline view across all spend regardless of source system.
 - As a Sourcing Associate, I want spend records enriched with Coupa/EAA detail (supplier, contract, category) when available, so I get richer context without losing coverage from Oracle-only transactions.
 - As a Data/AI Engineer, I want a defined precedence rule (source system detail overrides Oracle AP where both exist) so spend isn't double-counted or contradicted across sources.
-

Epic 2: Multi-Signal Spend Categorization & Taxonomy

Description: Build an indirect spend category taxonomy (mirroring the direct-side family tree) using Coupa commodity hierarchy as the anchor, supplemented by other categorization signals, so spend can be reliably classified even where source-system categories are weak or missing.

Backed by: Commodity (commodities_v1), Item.record.commodity.id , SpendFact , DataQualityIssue

Signals identified in session (in priority order to test):

1. Coupa category field / commodity hierarchy (source of truth for MVP)
2. GL line-item description (via distribution table - need Nick/Matt's team to expose this; only payments table currently shared)
3. OCR-extracted invoice line-item detail
4. Supplier Hub category/subcategory + D&B codes

User Stories

- As a Data/AI Engineer, I want spend categorized against the Coupa commodity hierarchy as the default taxonomy, so category is available consistently across all Coupa-sourced spend.
 - As a Sourcing Associate, I want spend that can't be confidently categorized to be flagged rather than silently miscategorized, so I know which numbers to trust.
 - Degree of confidence can be updated, not hard-coded
 - As a Data/AI Engineer, I want access to the GL distribution table (not just the payments table), Supplier Hub category/subcategory and NAICS codes for a supplier, so line-item description can be used as a categorization signal.
 - As a Sourcing Associate, I want a simple way to confirm or correct a spend record's category, so the tool learns from my input and improves future categorization accuracy.
-

Epic 3: Persona-Based Access and Ontology-Style Drill-Through

Description: Deliver a hierarchical, drill-through spend view modeled on Palantir's ontology approach (interconnected objects, not just a flat report) so users can navigate from spend totals down to supplier, contract, and invoice-level detail.

Backed by: `SupplierProfile`, `ContractCoverage`, `LineItemEvidence`, `gold_object_relationship`

Confirmed persona/access model:

- Single shared view for MVP - every user (Sourcing Associate, COE, Executive Sponsor) sees the same unrestricted spend population; nobody in Indirect Sourcing is blocked from seeing everything
- Users manually filter and drill up/down to the level relevant to them (e.g. spend → supplier → contract → invoice); no automated persona-based pre-filtering at MVP
- Role-based permissions and role-specific default views (e.g. a CPO landing view vs. an individual associate landing view) are explicitly post-MVP

- Misty and Joe engaging core users for testing - building a plan for who reviews and provides input across all capabilities, not just this one

User Stories

- As a Sourcing Associate, I want to see the full spend population by default and manually filter down to my categories, so I can focus on what's relevant without waiting for automated persona-based filtering.
 - As any user (Sourcing Associate, COE, or Executive Sponsor), I want to navigate from a spend total to the underlying supplier, then to that supplier's contracts, then to related invoices, so I can investigate without re-querying from scratch.
-

Epic 4: Managed vs. Unmanaged Spend Visibility

Description: Distinguish spend covered by an active contract (managed) from spend outside contract coverage (unmanaged), as one of the four core drill-down dimensions for Spend Visibility.

Backed by: `ContractCoverage`, `contract_coverage` metric, `gold_contract_coverage_metric`

Confirmed drill-down/up dimensions (MVP): category, supplier, managed vs. unmanaged spend, location

User Stories

- As a Sourcing Associate, I want to see what percentage of my category's spend is covered by an active contract, so I can identify leakage or renegotiation opportunities.
 - As a Sourcing Associate, I want to drill from a managed/unmanaged split down to the specific suppliers and transactions driving unmanaged spend, so I can act on it.
 - As a Sourcing Associate, I want to filter the spend view by location/division in addition to category and supplier, so I can see how spend patterns vary geographically.
-

Epic 5: Data Quality and Trust Signals

Description: Surface freshness, parsing confidence, and known data-quality issues alongside spend figures, so users can gauge how much to trust a given number rather than treating all data as equally reliable.

Backed by: `DataQualityIssue`, `gold_source_freshness_metric`, `stale_data_age`

User Stories

- As a Sourcing Associate, I want to see how recently a spend figure was refreshed, so I know if I'm looking at current or stale data.
- As a Sourcing Associate, I want spend records with data-quality issues (missing PO, missing supplier, broken relationship) flagged, so I don't act on unreliable numbers.
- As a Data/AI Engineer, I want data-quality issues to block or qualify downstream AI recommendations (not just be logged), so bad data doesn't silently propagate into sourcing suggestions.