

User Story 2.4: Categorization Feedback Loop

As a Sourcing Associate, I want a simple way to confirm or correct a spend record's category, so the tool learns from my input and improves future categorization accuracy.

Description:

Provide a lightweight interface for users to mark a categorized spend record as correct or incorrect, and to submit a corrected category when needed - capturing that feedback as structured input to refine the categorization model over time.

Narrative:

No categorization approach - whether Coupa, the new business taxonomy, or additional signals - will be perfect on day one. Rather than treating categorization as a static, one-time assignment, users need a fast way to flag errors as they encounter them in the normal course of their work. This mirrors the category-validation approach Palantir uses (accept/reject/recategorize, feeding retraining) and the same feedback pattern already planned for Sourcing Automation recommendations - categorization should follow the same principle.

Business Value:

Turns every user interaction into a potential improvement to the model, rather than requiring a separate, disconnected review exercise. Directly supports the iterative retraining cycles already planned for the taxonomy rollout (per the capacity planning model), and reduces the volume of low-confidence flags (Story 2.2) over time as the model learns from real corrections.

Acceptance Criteria:

- Given a spend record's category, when a user reviews it, then they can confirm it's correct or flag it as incorrect with minimal friction (e.g. one click)
- Given a user flags a category as incorrect, when doing so, then they can specify the correct category
- Given feedback has been submitted, when captured, then it is stored as structured data (not just logged), available to inform model retraining

Out of Scope:

- Automatically retraining the model in real time based on a single correction - this story covers capturing feedback, not the retraining cadence/mechanism itself
- Bulk/batch correction tools - this story covers individual record-level feedback only