

User Story 5.3: Data Quality Gating for AI Recommendations

As a Data/AI Engineer, I want data-quality issues to block or qualify downstream AI recommendations (not just be logged), so bad data doesn't silently propagate into sourcing suggestions.

Description:

Ensure that data-quality issues identified in spend records (per Story 5.2) actively influence downstream AI-generated recommendations - either blocking a recommendation from being generated, or qualifying it with a visible caveat - rather than being recorded only as a passive log entry with no effect on what gets surfaced to users.

Narrative:

Flagging a data-quality issue (Story 5.2) only protects users who are actively looking at that specific record. But Sourcing Automation and Invoice Leakage generate recommendations that aggregate and interpret spend data - if the underlying data feeding those recommendations is unreliable, the recommendation itself inherits that unreliability, potentially without the user ever realizing the root data was flawed. This story ensures data-quality signals actively gate or qualify what gets recommended, not just what gets displayed.

Business Value:

Protects the credibility of Sourcing Automation and Invoice Leakage's recommendations - a single bad recommendation built on flawed data can undermine trust in the entire recommendation engine, even if most recommendations are sound. This is foundational to both of those capabilities functioning as "accurate, actionable" per their stated success criteria.

Acceptance Criteria:

- Given a spend record feeding into a potential recommendation has a flagged data-quality issue, when the recommendation would otherwise be generated, then it is either suppressed or generated with a visible data-quality caveat attached
- Given a recommendation is qualified with a data-quality caveat, when a user views it, then the specific underlying issue is identifiable, not just a generic disclaimer
- Given a data-quality issue is later resolved, when recommendations are regenerated, then previously blocked/qualified recommendations are re-evaluated without the caveat

Out of Scope:

- The specific recommendation logic itself (non-catalog match, contract leakage, etc.) - covered under Sourcing Automation's own epics; this story only covers the data-quality gate

applied to that logic

- Automatically resolving the underlying data-quality issue - covered by Story 5.2's flagging, not this story's remediation